

# Information Retrieval

CS 547/DS 547

Worcester Polytechnic Institute

Department of Computer Science

Instructor: Prof. Kyumin Lee

# Project

- 3~4 person team
- Notify names of your project team members by February 2
- Dates:
  - [7%] Project Proposal Writing: March 17 by 11:59pm
  - [5%] Project Proposal Presentation: March 21
  - [8%] Project website: April 25 by 11:59pm
  - [11%] Project Workshop: April 26 in-class
- [https://canvas.wpi.edu/courses/46542/pages/project?module\\_item\\_id=888446](https://canvas.wpi.edu/courses/46542/pages/project?module_item_id=888446)

# Previous Year's Projects

- <https://exquisite-chebakia-9a513b.netlify.app/>
- <https://newsinspector.github.io/>
- <https://sites.google.com/view/newsbaordrecommender/home>
- <https://kratikashetty.github.io/CS547-Information-Retrieval/>
- <https://wheeleddoors.github.io/index/>
- <https://yelp-recommendation.ue.r.appspot.com/report>
- <https://github.com/khordoo/disaster-watch-classifier>

# Previous Class...

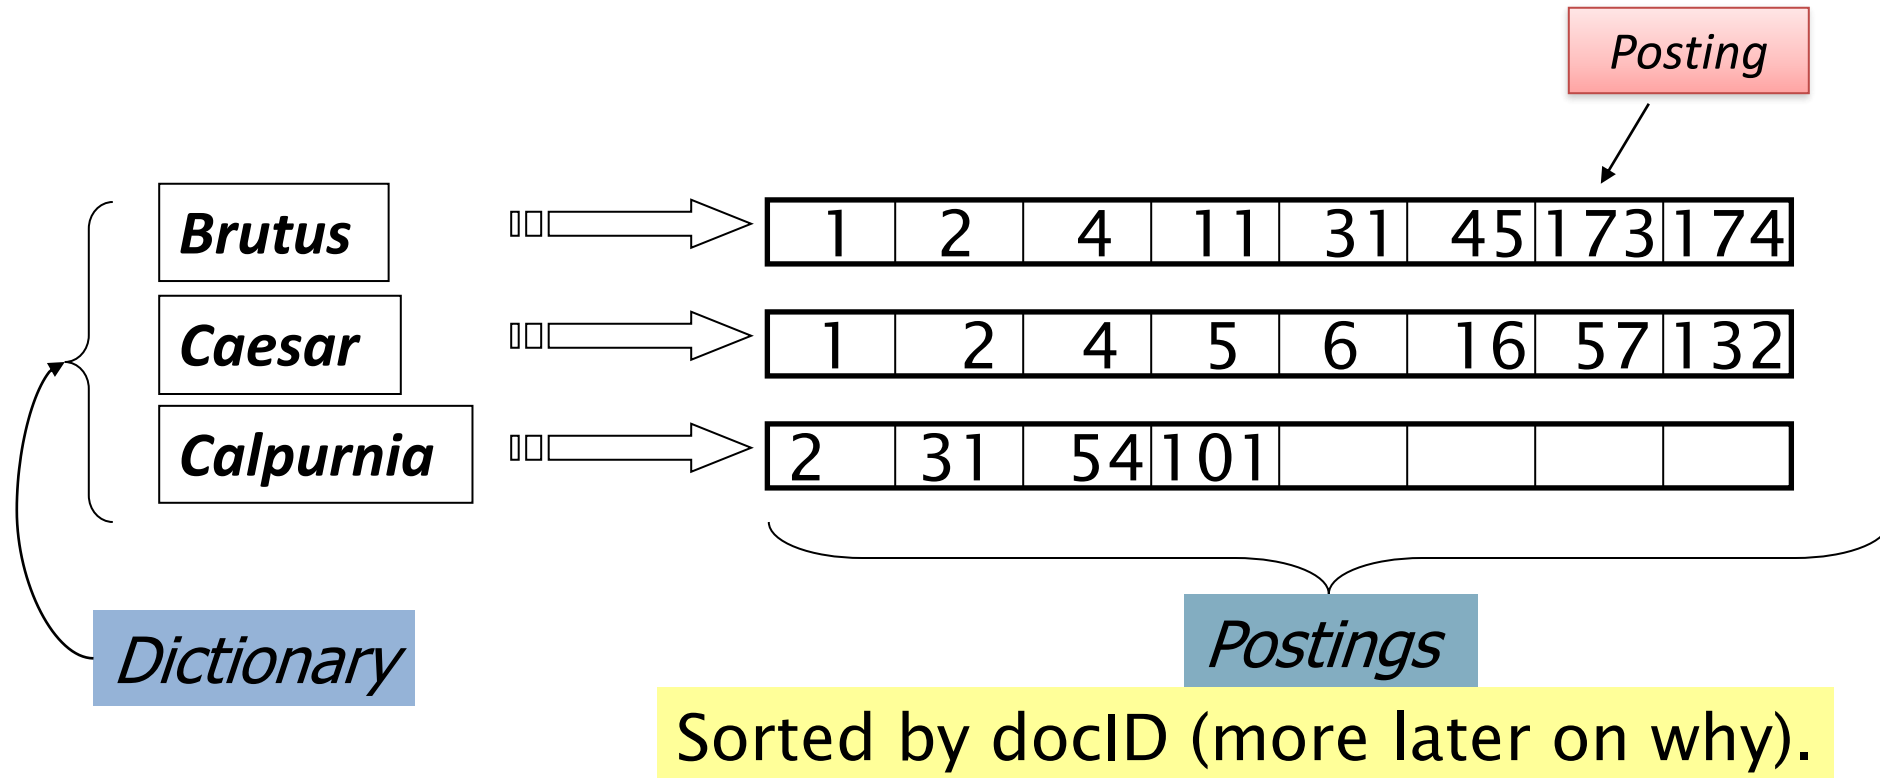
Boolean Retrieval  
Model

# Previous Class...

Boolean Retrieval  
Model

Inverted index

# Inverted index




# Previous Class...

Query Optimization

# Query optimization

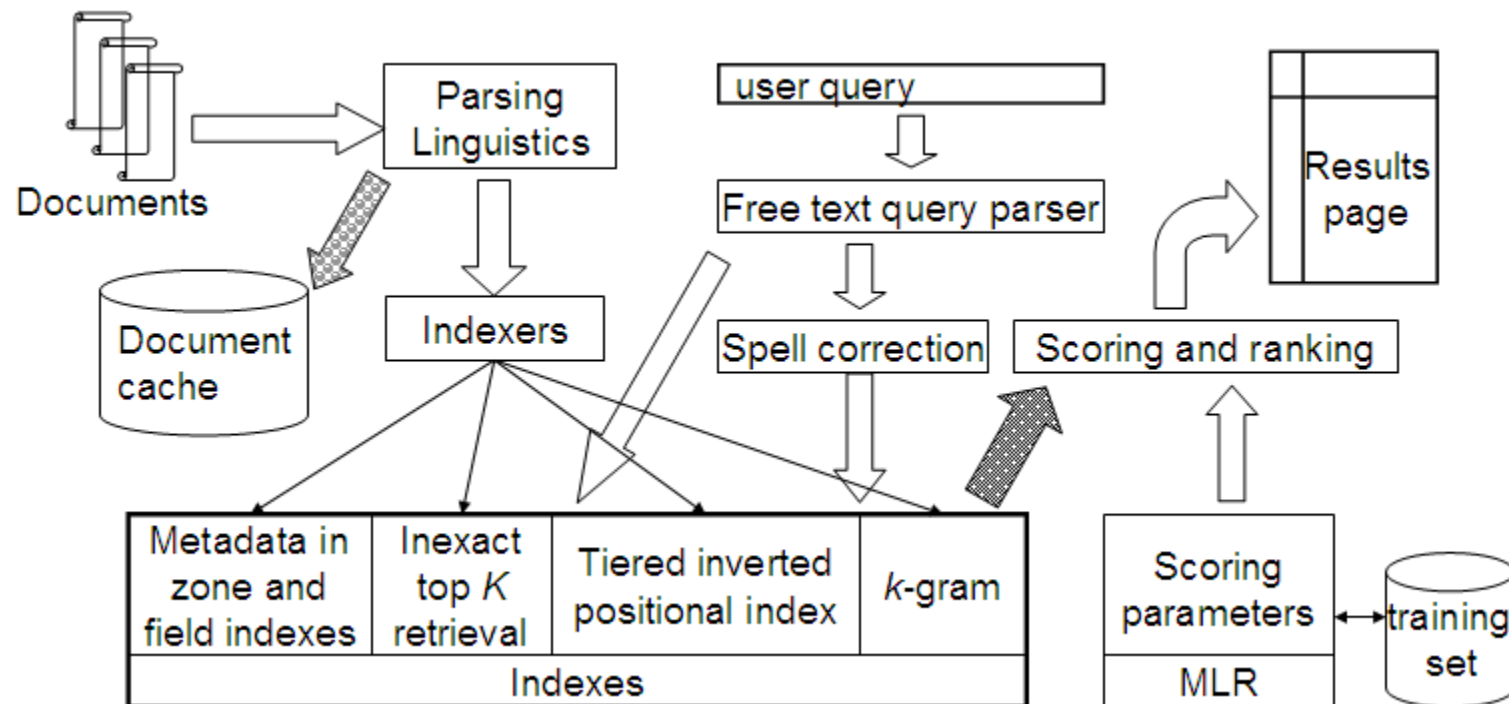
- Consider a query that is an *AND* of  $n$  terms.
- For each of the  $n$  terms, get its postings, then *AND* them together.
- What is the best order for query processing?



Query: **Brutus AND Calpurnia AND Caesar**



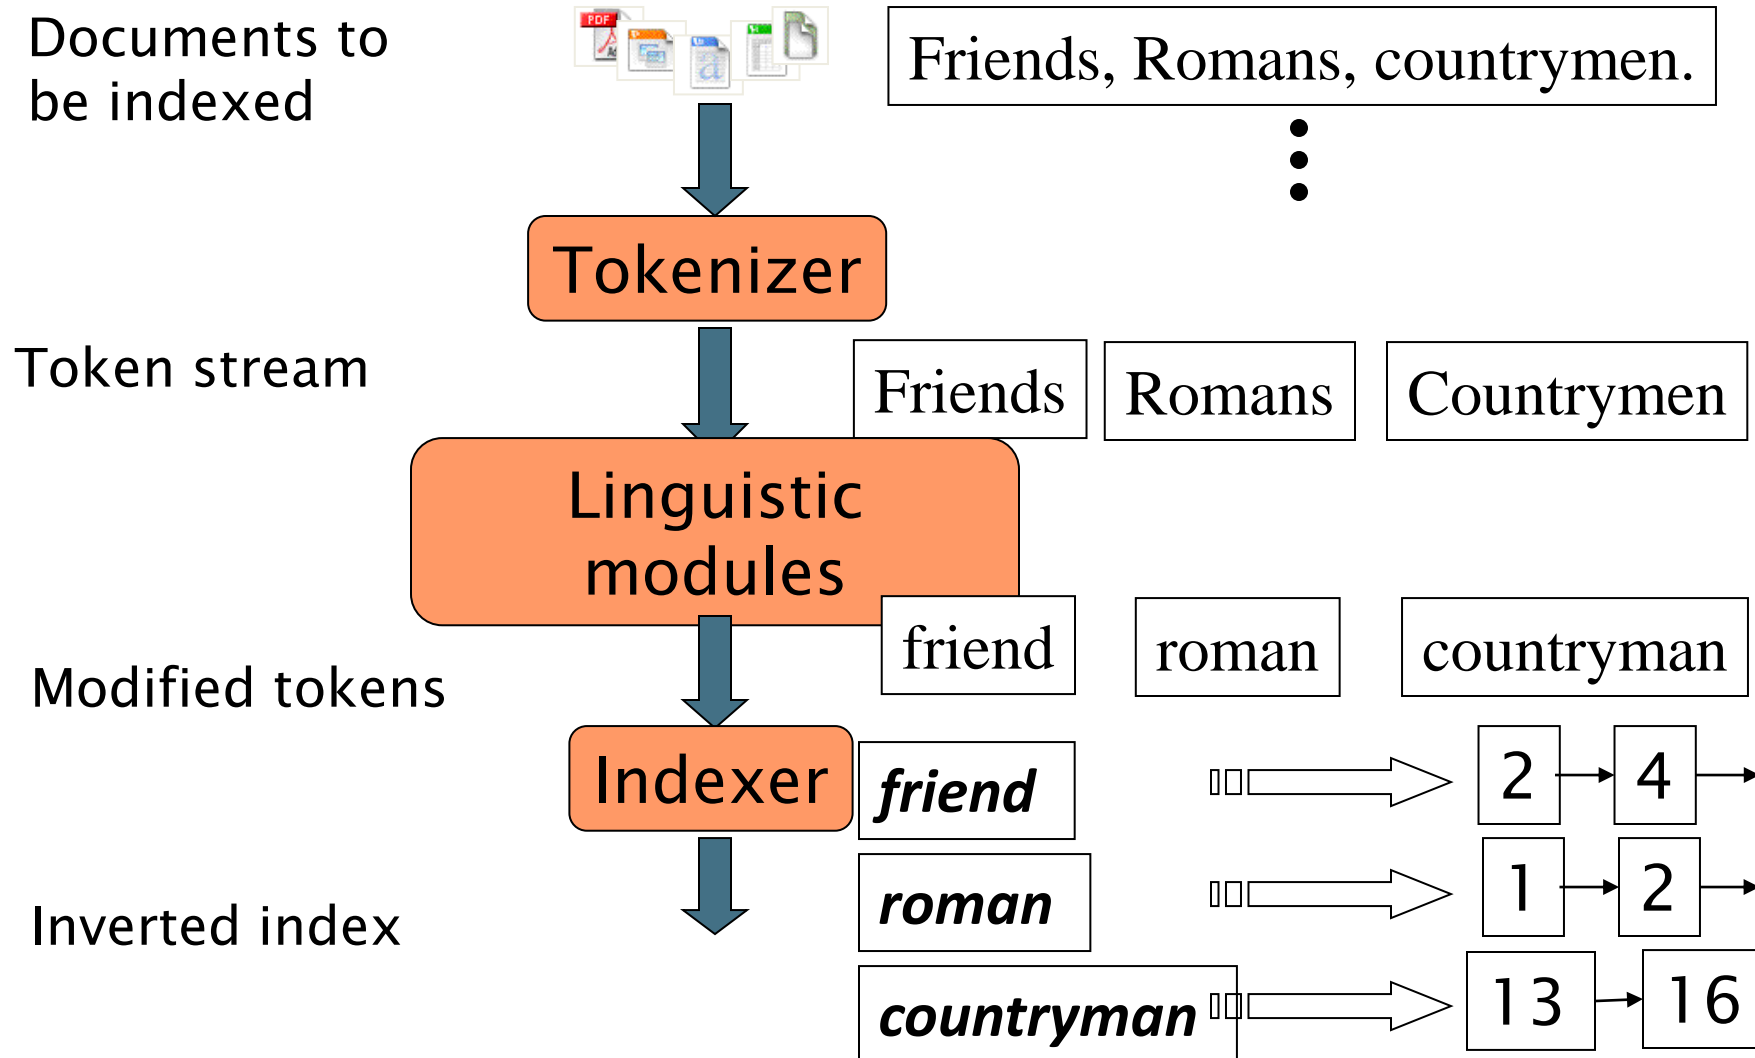
# What's next ...



# Our assumptions so far

- We know what a document is
- We know what a term is
  - In reality, it can be complex
- So... We'll look at how we define and process the vocabulary of terms in a collection

# Recall the basic indexing pipeline



# Initial stages of text processing

---

- Tokenization
  - Cut character sequence into word tokens
    - Deal with ***“John’s”, a state-of-the-art solution***
- Normalization
  - Map text and query term to same form
    - You want **U.S.A.** and **USA** to match
- Stemming
  - We may wish different forms of a root to match
    - ***authorize, authorization***
- Stop words
  - We may omit very common words (or not)
    - ***the, a, to, of***

# Parsing a document

---

- What format is it in?
  - pdf/word/excel/html?
- What language is it in?
- What character set is in use?
  - (CP1252, UTF-8, ...)

Each of these is a classification problem, which we will study later in the course.

But these tasks are often done heuristically ...

# Tokenization

---

- Input: “*Friends, Romans and Countrymen*”
- Output: Tokens
  - *Friends*
  - *Romans*
  - *Countrymen*
- A **token** is an instance of a sequence of characters
- Each such token is now a candidate for an index entry, after further processing
  - Described below
- But what are valid tokens to emit?

# Why tokenization is difficult -- even in English

---

- Example: *Mr. O'Neill thinks that the boys' stories about Chile's capital aren't amusing.*
- **Tokenize this sentence**

# One word or two? (or several)

---

- Hewlett-Packard
- State-of-the-art
- co-education
- the hold-him-back-and-drag-him-away maneuver
- data base
- San Francisco
- Los Angeles-based company
- cheap San Francisco-Los Angeles fares
- York University vs. New York University



# Numbers

---

- 3/12/91
- 12/3/91
- Mar 12, 1991
- B-52
- 100.2.86.144
- (800) 234-2333
- 800.234.2333

# Chinese: No whitespace

---

莎拉波娃现在居住在美国东南部的佛罗里达。今年4月9日，莎拉波娃在美国第一大城市纽约度过了18岁生日。生日派对上，莎拉波娃露出了甜美的微笑。

# Bidirectionality in Arabic

---

- Arabic (or Hebrew) is basically written right to left, but with certain items like numbers written left to right
- Words are separated, but letter forms within a word form complex ligatures

استقلت الجزائر في سنة 1962 بعد 132 عام من الاحتلال الفرنسي.

- ‘Algeria achieved its independence in 1962 after 132 years of French occupation.’

# Normalization

---

- Need to “normalize” words in indexed text as well as query words into the same form
  - We want to match ***U.S.A.*** and ***USA***
- We most commonly implicitly define **equivalence classes** of terms
  - e.g., deleting periods to form a term
- Alternative is to do asymmetric expansion:
  - Enter: *window*            Search: *window, windows*
  - Enter: *windows*           Search: *Windows, windows*
  - Enter: *Windows*           Search: *Windows*
- Potentially more powerful, but less efficient

# Case folding

---

- Reduce all letters to lower case
  - exception: upper case in mid-sentence?
    - e.g., General Motors
    - Fed vs. fed
    - SAIL vs. sail
- Often best to lower case everything, since users will use lowercase regardless of 'correct' capitalization...

# Stop words

---

- With a stop list, you exclude from the dictionary entirely the commonest words. Intuition:
  - They have little semantic content: *the, a, and, to, be*
  - There are a lot of them: ~30% of postings for top 30 words
- But the trend is away from doing this:
  - Good compression techniques means the space for including stop words in a system is very small
  - Good query optimization techniques mean you pay little at query time for including stop words.
  - You need them for:
    - Phrase queries: “King of Denmark”
    - Various song titles, etc.: “Let it be”, “To be or not to be”
    - “Relational” queries: “flights to London”

# Lemmatization

---

- Reduce inflectional/variant forms to base form
- Example: *am, are, is* → *be*
- Example: *car, cars, car's, cars'* → *car*
- Example: the boy's cars are different colors → the boy car be different color
- Lemmatization implies doing “proper” reduction to dictionary headword form (the lemma).

# Stemming

---

- Reduce terms to their “roots” before indexing
- “Stemming” suggests crude affix chopping
- language dependent
- Example: ***automate(s), automatic, automation*** all reduced to ***automat***.



# Porter Stemming Algorithm

---

- Most common algorithm for stemming English
- Results suggest that it is at least as good as other stemming options
- Contains 5 phases of reductions
- Phases are applied sequentially
- Each phase consists of a set of commands.
  - Sample command: Delete final element if what remains is longer than 1 character
  - replacement → replac
  - cement → cement

# Porter stemmer:

## A few rules

---

### Rule

SSSES → SS

IES → I

SS → SS

S →

### Example

caresses → caress

ponies → poni

caress → caress

cats → cat

# Three stemmers: A comparison

---

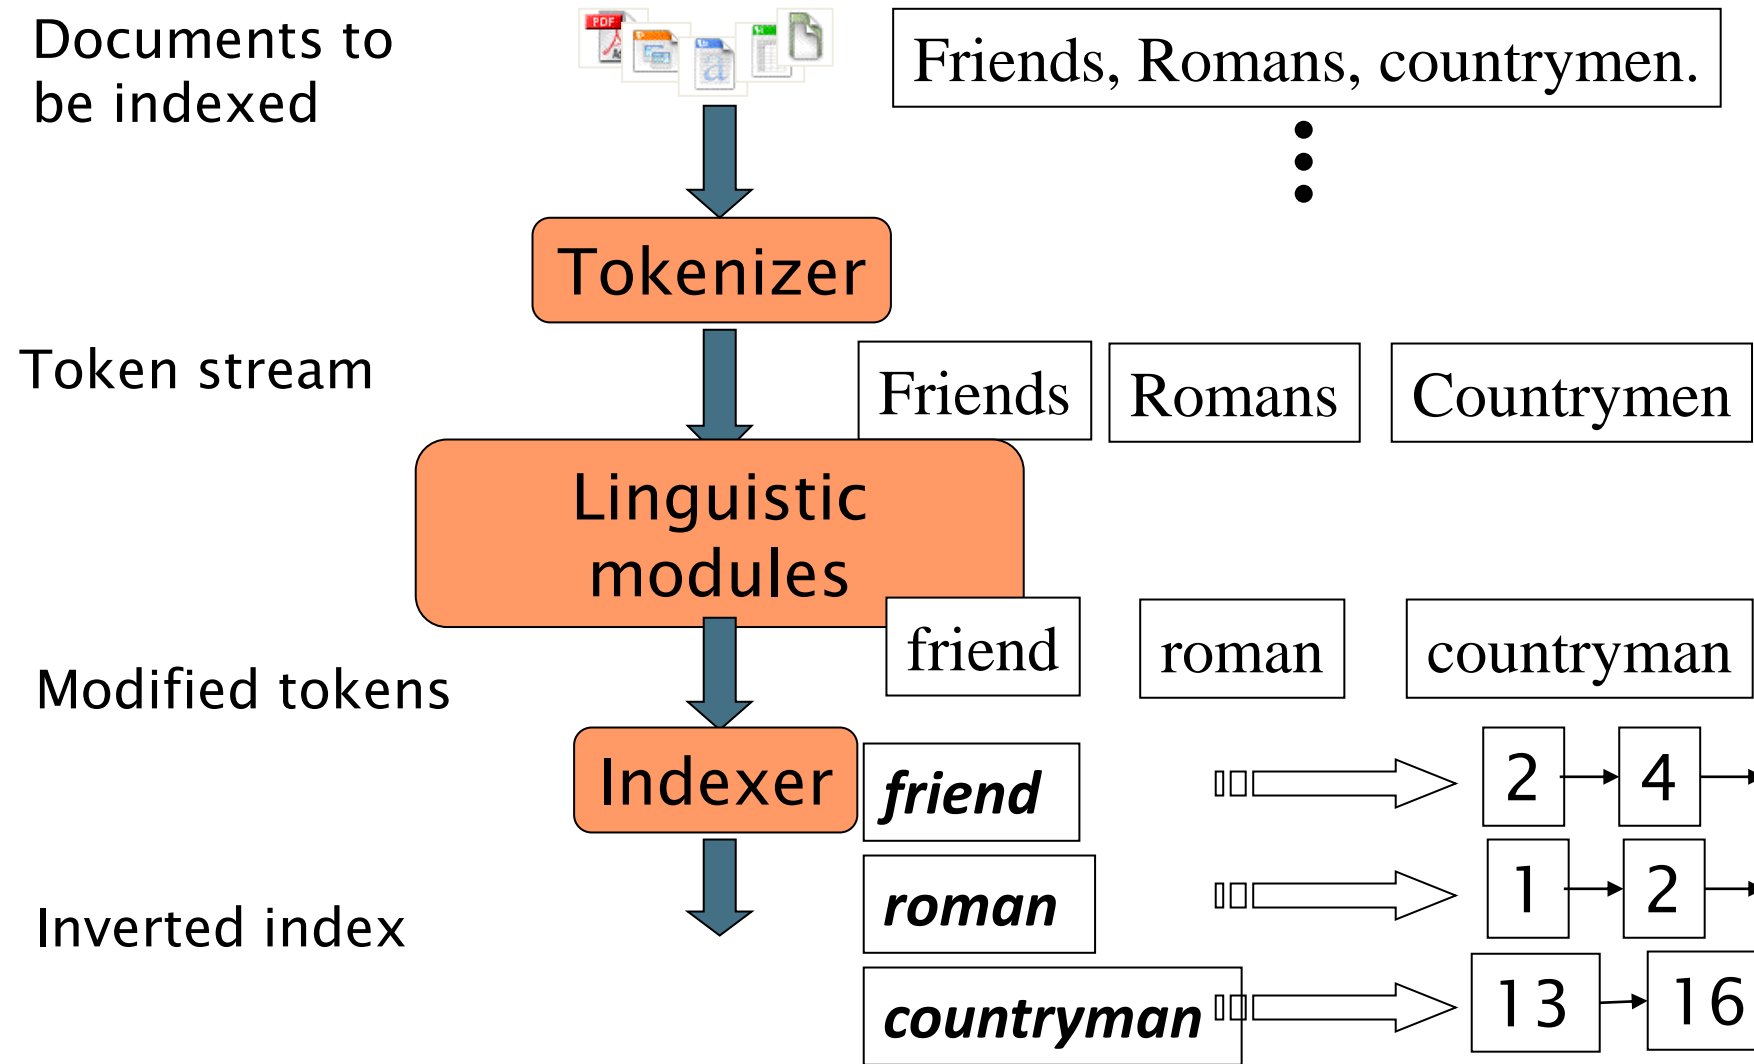
- **Sample text:** Such an analysis can reveal features that are not easily visible from the variations in the individual genes and can lead to a picture of expression that is more biologically transparent and accessible to interpretation
- **Porter stemmer:** such an analysis can reveal features that are not easily visible from the variations in the individual genes and can lead to a picture of expression that is more biologically transparent and accessible to interpretation
- **Lovins stemmer:** such an analysis can reveal features that are not easily visible from the variations in the individual genes and can lead to a picture of expression that is more biologically transparent and accessible to interpretation
- **Paice stemmer:** such an analysis can reveal features that are not easily visible from the variations in the individual genes and can lead to a picture of expression that is more biologically transparent and accessible to interpretation

# Does stemming improve effectiveness?

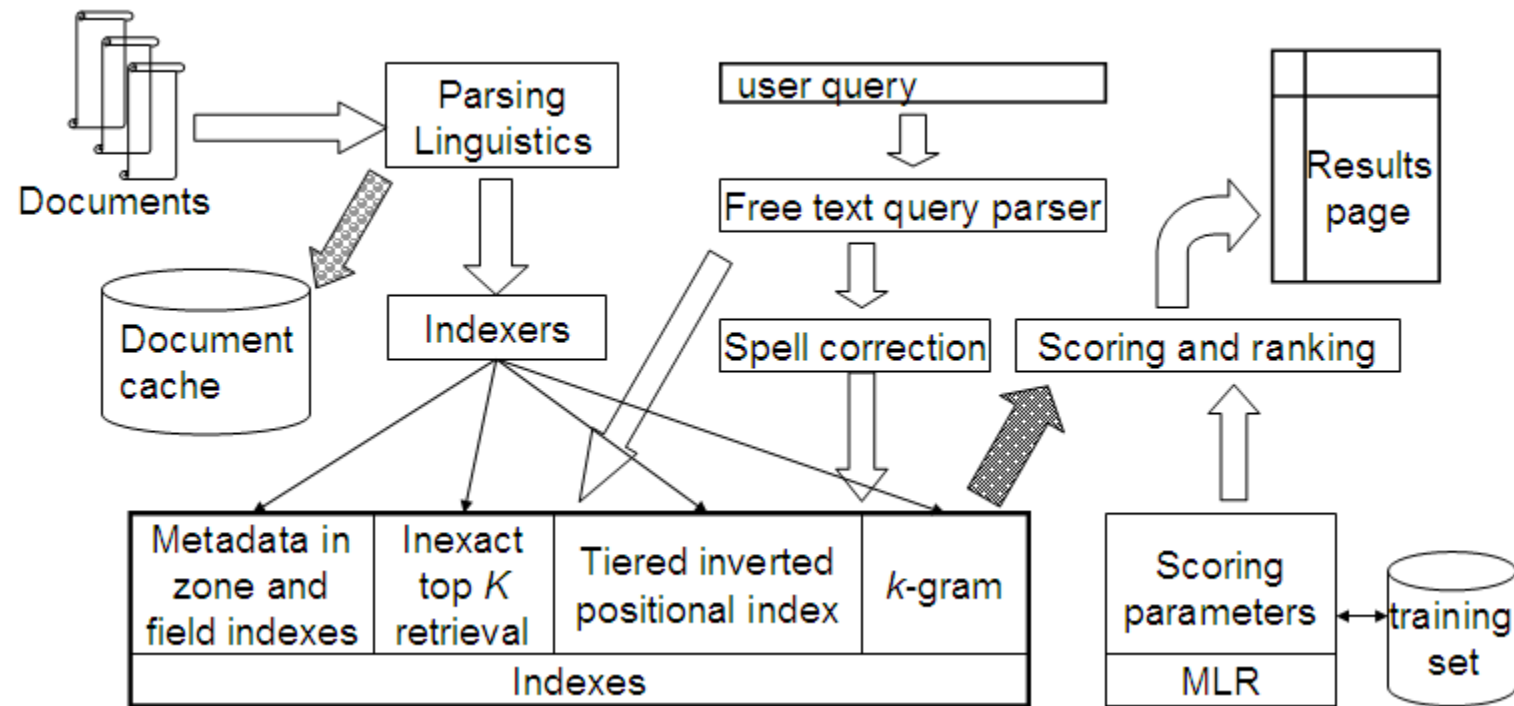
---

- In general, stemming increases effectiveness for some queries, and decreases effectiveness for others.
- Porter Stemmer equivalence class **oper** contains all of **operate operating operates operation operative operatives operational**.
- Queries where stemming hurts: “operational AND research”, “operating AND system”, “operative AND dentistry”

# Recall the basic indexing pipeline



# Big Picture



# HW1

---

- [https://canvas.wpi.edu/courses/46542/assignments/283401?module\\_item\\_id=888447](https://canvas.wpi.edu/courses/46542/assignments/283401?module_item_id=888447)

# Next...

---

- Need a better index than simple <term: docs>
- How can we improve on our basic index?
  - **Skip pointers:** faster postings merges
  - **Positional index:** Phrase queries and Proximity queries
  - **Permuterm index:** Wildcard queries




---

Faster postings merges:  
Skip pointer

# Recall basic merge

- Walk through the two postings simultaneously, in time linear in the total number of postings entries





If the list lengths are  $m$  and  $n$ , the merge takes  $O(m+n)$  operations.

Can we do better?

Yes (if the index isn't changing too fast).


# Augment postings with skip pointers (at indexing time)

---



- Why?
- To skip postings that will not figure in the search results.
- How?
- Where do we place skip pointers?

# Query processing with skip pointers



Suppose we've stepped through the lists until we process **8** on each list. We match it and advance.

We then have **41** and **11** on the lower. **11** is smaller.


But the skip successor of **11** on the lower list is **31**, so we can skip ahead past the intervening postings.

INTERSECTWITHSKIPS( $p_1, p_2$ )

```
1  answer  $\leftarrow \langle \rangle$ 
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
3  do if  $\text{docID}(p_1) = \text{docID}(p_2)$ 
4      then  $\text{ADD}(\text{answer}, \text{docID}(p_1))$ 
5           $p_1 \leftarrow \text{next}(p_1)$ 
6           $p_2 \leftarrow \text{next}(p_2)$ 
7  else if  $\text{docID}(p_1) < \text{docID}(p_2)$ 
8      then if  $\text{hasSkip}(p_1)$  and  $(\text{docID}(\text{skip}(p_1)) \leq \text{docID}(p_2))$ 
9          then while  $\text{hasSkip}(p_1)$  and  $(\text{docID}(\text{skip}(p_1)) \leq \text{docID}(p_2))$ 
10             do  $p_1 \leftarrow \text{skip}(p_1)$ 
11             else  $p_1 \leftarrow \text{next}(p_1)$ 
12      else if  $\text{hasSkip}(p_2)$  and  $(\text{docID}(\text{skip}(p_2)) \leq \text{docID}(p_1))$ 
13          then while  $\text{hasSkip}(p_2)$  and  $(\text{docID}(\text{skip}(p_2)) \leq \text{docID}(p_1))$ 
14             do  $p_2 \leftarrow \text{skip}(p_2)$ 
15             else  $p_2 \leftarrow \text{next}(p_2)$ 
16  return answer
```

# Where do we place skips?

- Tradeoff:
  - More skips  $\rightarrow$  shorter skip spans  $\Rightarrow$  more likely to skip. But lots of comparisons to skip pointers.
  - Fewer skips  $\rightarrow$  few pointer comparison, but then long skip spans  $\Rightarrow$  few successful skips.



# Placing skips

---

- So... More skips or fewer skips... Where to add skip pointers???
- Simple heuristic: for postings of length  $L$ , use  $\sqrt{L}$  evenly-spaced skip pointers
- Easy if the index is relatively static; harder if  $L$  keeps changing because of updates.

---

# Positional Index



# Phrase queries

---

- Want to be able to answer queries such as “***stanford university***” – as a phrase
- Thus the sentence “*I went to university at Stanford*” is not a match.
  - The concept of phrase queries has proven easily understood by users; about 10% of web queries are phrase queries
- How??

# A first attempt: Biword indexes

---

- Index every consecutive pair of terms in the text as a phrase
- For example the text “Friends, Romans, Countrymen” would generate the biwords
  - *friends romans*
  - *romans countrymen*
- Each of these biwords is now a dictionary term
- Two-word phrase query-processing is now immediate.

# Longer phrase queries

---

- Longer phrases can be processed by breaking them down?
- ***stanford university palo alto*** can be broken into the Boolean query on biwords:

***stanford university AND university palo  
AND palo alto***

Any problem?

Without the docs, we cannot verify that the docs matching the above Boolean query do contain the phrase.



Can have false positives!

## Solution 2: Positional indexes

---

- In the postings, store, for each ***term*** the position(s) in which tokens of it appear:

<***term***, number of docs containing ***term***;

*doc1*: position1, position2 ... ;

*doc2*: position1, position2 ... ;

etc.>

# Positional index example

---

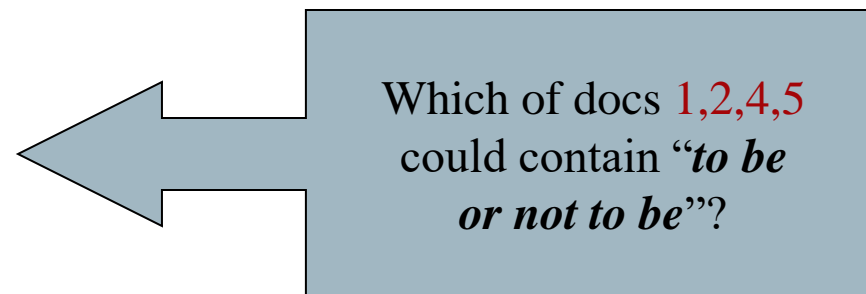
<*be*: 993427;

*1*: 7, 18, 33, 72, 86, 231;

*2*: 3, 149;

*4*: 17, 191, 291, 430, 434;

*5*: 363, 367, ...>



- Can compress position values/offsets
- Nevertheless, this expands postings storage *substantially*

# Processing a phrase query

---

- Extract inverted index entries for each distinct term: ***to, be, or, not.***
- Merge their *doc:position* lists to enumerate all positions with “***to be or not to be***”.
  - ***to:***
    - 2:1,17,74,222,551; 4:8,16,190,429,433; 7:13,23,191; ...
  - ***be:***
    - 1:17,19; 4:17,191,291,430,434; 5:14,19,101; ...
- Same general method for proximity searches

# Proximity queries

---

- Employment /3 place
  - Here, / $k$  means “within  $k$  words of (on either side)”.
- Clearly, positional indexes can be used for such queries; biword indexes cannot.

# Proximity Queries in Search Engines

---

- Google Search supports
  - keyword1 AROUND(n) keyword2
- Bing
  - keyword1 near:n keyword2 where n=the number of maximum separating words.
- Yahoo
  - keyword1 NEAR keyword2
- Exalead
  - keyword1 NEAR/n keyword2 where n is the number of words.

E.g., hotel around(5) terminal vs hotel around(3) terminal at Google



# Positional index size

- Need an entry for each occurrence, not just once per document
- Index size depends on average document size
  - Average web page has <1000 terms
  - SEC filings, books, even some epic poems ... easily 100,000 terms
- Consider a term with frequency 0.1%

Document size	Postings	Positional postings
1000	1	1
100,000	1	100

# Positional index size

---

- You can compress position values/offsets
- Nevertheless, a positional index expands postings storage substantially
- Nevertheless, it is now standardly used because of the power and usefulness of phrase and proximity queries ... whether used explicitly or implicitly in a ranking retrieval system.

# Rules of thumb

---

- A positional index is 2–4 as large as a non-positional index
- Positional index size 35–50% of volume of original text

# Positional Indexes: Wrap-up

---

- With a positional index, we can answer
  - phrase queries
  - proximity queries

# Today...

- Need a better index than simple <term: docs>
- How can we improve on our basic index?
  - **Skip pointers**: faster postings merges
  - **Positional index**: Phrase queries and Proximity queries
  - **Permuterm index**: Wildcard queries

Wild-card queries

# Wild-card queries: \*

---


- ***mon\****: find all docs containing any word beginning with “mon”.
- Easy with binary tree (or B-tree) lexicon: retrieve all words in range: ***mon***  $\leq w$  ***< moo***





# Tree: B-tree

---



- Definition: Every internal node has a number of children in the interval  $[a,b]$  where  $a, b$  are appropriate natural numbers, e.g.,  $[2,4]$ .

# Wild-card queries: \*

---

- ***mon\****: find all docs containing any word beginning with “mon”.
- Easy with binary tree (or B-tree) lexicon: retrieve all words in range: ***mon***  $\leq w$  ***< moo***
- ***\*mon***: find words ending in “mon”: harder
  - Maintain an additional B-tree for terms *backwards*.  
Can retrieve all words in range: ***nom***  $\leq w$  ***< non***.

Exercise: from this, how can we enumerate all terms meeting the wild-card query ***pro\*cent*** ?

# Query processing

---

- At this point, we have an enumeration of all terms in the dictionary that match the wild-card query.
- We still have to look up the postings for each enumerated term.
- E.g., consider the query:

***se\*ate AND fil\*er***

This may result in the execution of many Boolean *AND* queries.


# B-trees handle \*'s at the end of a query term

---


- How can we handle \*'s in the middle of query term?
  - *co\*tion*
- We could look up *co\** AND *\*tion* in a B-tree and intersect the two term sets
  - Expensive
- The solution: transform wild-card queries so that the \*'s occur at the end
- This gives rise to the **Permuterm** Index.

# Permuterm index

- For term ***hello***, index under:
  - ***hello\$, ello\$h, llo\$he, lo\$hel, o\$hell, \$hello***
 where \$ is a special symbol.



- Queries:
  - **X** lookup on **X\$**      **X\*** lookup on **\$X\***
  - **\*X** lookup on **X\$\***      **\*X\*** lookup on **X\***
  - **X\*Y** lookup on **Y\$X\***
  - **X\*Y\*Z** ??? Exercise!



# Permuterm query processing

---

- Rotate query wild-card to the right
- Now use B-tree lookup as before.
- *Permuterm problem:  $\approx$  quadruples lexicon size*



Empirical observation for English.

---

Any Questions?